

PREM TIWARI

Software Engineer II · Backend & Data Engineering

Bengaluru, Karnataka, India • +91 96448 46577 • sepremtiwari@gmail.com

Portfolio • [linkedin.com/in/premtiwari14f](https://www.linkedin.com/in/premtiwari14f) • github.com/seprem • leetcode.com/u/prem14f

PROFESSIONAL SUMMARY

Software Engineer II based in **Bengaluru**, with **3 years 8 months** of experience moving from **wind-farm IoT** to large-scale **data platforms**. At Vestas I built **real-time IoT** ingestion, including an **OPC-UA** pipeline from scratch on **Kafka**, **MQTT**, and **Kinesis**. At Smartsheet I deliver **config-driven Databricks ingestion** and **CDC** across **US**, **EU**, and **AU** for sources including **MySQL CDC**, **DynamoDB**, **Kinesis**, **Salesforce API**, **NetSuite**, and **Stripe**, plus an **OpenSearch** pipeline. Tech I work on: **Python**, **Go**, **SQL**, **PySpark**, **Kafka**, **Flink**, **OPC-UA**, **MQTT**, **Databricks**, **Snowflake**, **Airflow**, and **AWS**.

PROFESSIONAL EXPERIENCE

Software Engineer II — *Smartsheet*

Oct 2025 – Present

- Built **config-driven ingestion for different sources** on **Databricks** across **US**, **EU**, and **AU**: **1,045 jobs / 2,756 tasks** over **11 sources** (**MySQL CDC**, **DynamoDB**, **Kinesis**, **Salesforce API**, **NetSuite**, **Stripe**), plus a **Snowflake → Delta** backfill (**42 jobs / 287 tasks**), shipped with **Databricks Asset Bundles (CI/CD)**.
- Migrated **multi-task Spark streaming jobs** on **Databricks** from **all-purpose** to **job clusters** across **US**, **EU**, and **AU** production, cutting compute costs by **\$200K+ annually**.
- Landed CDC from **82 sharded Aurora PostgreSQL** databases (**Aurora PostgreSQL → Flink → Kinesis → Delta tables**): **876 jobs / 1,547 tasks**, **282 streaming**, feeding a **vector database for AI and LLMs**.
- Built a **high-availability log pipeline** (**CloudWatch → subscription filter → Kinesis → Logstash → OpenSearch**) ingesting **45–50 TB/day** across **50–100 Kinesis streams** and **1,000+ shards**, with **sub-second search**.
- Built **audit automation** for **Databricks**, **Snowflake**, **Jenkins**, and **AWS**, plus **onboarding docs** and an **access-request dashboard** so other teams **self-serve** access — cutting **weeks of manual effort to days or hours**. Drove **on-call tickets** from **~40 to single digits** and **daily ingestion failures** by **~34%**.

Data Engineer — *Vestas*

Feb 2024 – Sep 2025

- Designed a **real-time alarm platform** for **~150–200 farms / 19k turbines** on **Kafka**, **Redis**, and **Kinesis (Docker)** — **Protobuf-encoded alarms from Kafka**; built the **frontend API** serving turbine status from **DynamoDB** with **1–2s lag**, plus a **cursor-paginated file-download API** (**DocumentDB + S3**); streamed alarms via **Kinesis into Delta** for analytics.
- Built an **OPC-UA IoT pipeline from scratch** after an **AWS Greengrass POC** fell short — **22 Kafka partitions** ingesting **~10k data points/sec**, with **local-server buffering**; enhanced data availability by **96%**.
- Re-built a legacy **Databricks streaming job** as a **Dockerized service (MSK → Kinesis)** with **KEDA autoscaling** on **Kafka lag**, cutting **Delta ingestion latency** from minutes to **10–20 seconds**.
- Optimized **Databricks/PySpark** pipelines for solar & wind farm data, cutting processing time and latency by **~62%** by switching the source from **S3 to Kinesis**.

Assistant Data Engineer — *Vestas*

Feb 2023 – Jan 2024

- Built a **monitoring system** surfacing **data latency** across source services and pinpointing bottlenecks; reduced **API latency** by **36%**.
- Automated the previously manual, accuracy-based **ML model registration** to **AWS S3**, saving the team **12 hrs/week**; raised **unit-test coverage** to **95%**.
- Delivered **Databricks LTS** upgrades across **50+ jobs** with **data-quality testing** and **zero pipeline failures**.

TECHNICAL SKILLS

Languages: Python, Go, C++, SQL

Data Engineering: Spark, PySpark, Kafka, Flink, Databricks, Airflow, Airbyte, Snowflake, Delta Lake, DBT, ETL, Data Warehousing, Data Modeling, CDC

AWS & Cloud: EC2, S3, Lambda, Kinesis, MSK, SNS, SQS, Glue, API Gateway, VPC, CloudWatch, OpenSearch, Logstash

Databases: Aurora PostgreSQL, Aurora MySQL, PostgreSQL, MongoDB, DynamoDB, Redis

DevOps & Tools: Docker, Kubernetes, KEDA, Terraform, Terragrunt, Ansible, Jenkins, GitHub Actions, Grafana, MLflow

EDUCATION

B.Tech, Information Technology — Jabalpur Engineering College, Jabalpur, M.P.

Jul 2019 – Apr 2023

ACHIEVEMENTS

- **CodeChef:** 5-star, Global Rank 8, College Rank 1 (Starters 70)
- **LeetCode & GeeksforGeeks:** 450+ problems solved
- **TCS CodeVita:** Round 2 Global Rank 83
- **GFG Job-A-Thon 6:** Global Rank 76

RECOGNITION

- **Public kudos** at Smartsheet from **Engineering Manager** for **ownership**; recognized by **CTO, Chief AI Officer, SVP**, and **senior ICs** after leading stretch work equivalent to an **extra quarter** — including delivering in **one week** a project a contract team had been working on for **two months**.
- Won an internal **Smartsheet hackathon**: prototyped **usage analytics** that surface which product capabilities a customer needs from real usage — pitched as a **retention and expansion** play, reusable across similar accounts.

CERTIFICATIONS

- **Snowflake:** Collaboration & Marketplace, Data Application, Data Lake, Data Warehousing, Data Engineering